


# Scalable Kernel Density Classification via Threshold-Based Pruning

Edward Gan & Peter Bailis




# MacroBase: Analytics on Fast Streams

- Increasing Streaming Data
  - Manufacturing, Sensors, Mobile
  - Multi-dimensional + Latent anomalies
- Running in production
  - see CIDR17, SIGMOD17
- End-to-end operator cascades for:
  - Feature Transformation
  - Statistical Classification
  - Data Summarization




# Example: Space Shuttle Sensors



[UCI Repository]

8 Sensors Total



“Fuel Flow”




“Flight Speed”

Speed	Flow	Status
28	27	Fpv Close
34	43	High
52	30	Rad Flow
28	40	Rad Flow
...	...	

End-Goal: Explain anomalous speed / flow measurements.




Problem: Model distribution of speed / flow measurements.




# Difficulties in Data Modelling

Data Histogram




Gaussian Model




Poor Fit

# Difficulties in Data Modelling

Data Histogram




Mixture of Gaussians




Inaccurate: Gaps not captured

# Kernel Density Estimation (KDE)

Data Histogram



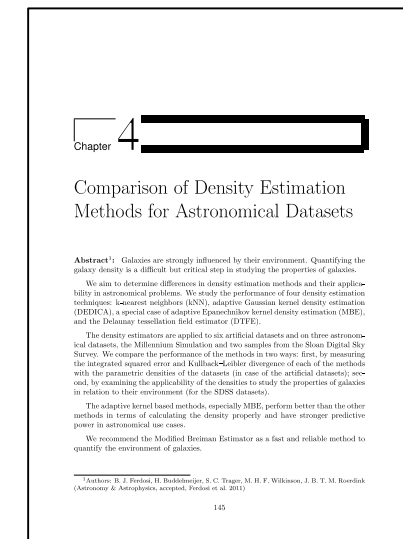



Kernel Density Estimate




Much better fit

# KDE: Statistical Gold Standard

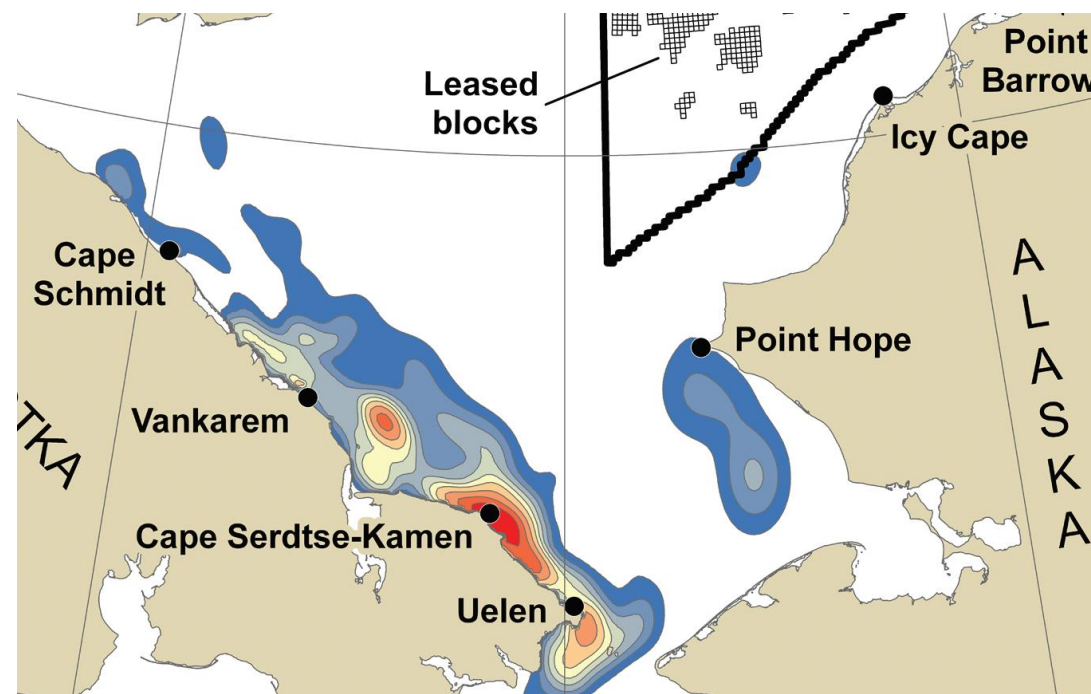
- Guaranteed to converge to the underlying distribution
- Provides normalized, true probability densities
- Few assumptions about shape of distribution: inferred from data



# KDE Usage



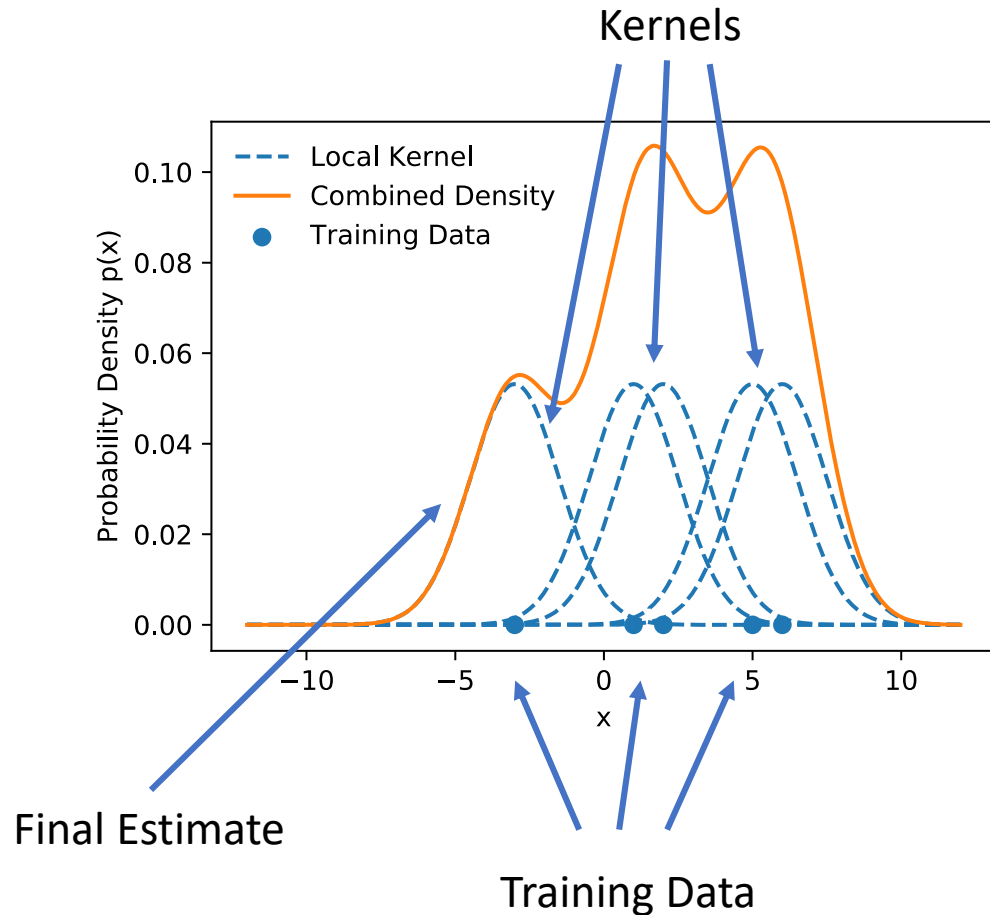
Galaxy Mass Distribution  
[Sloan Digital Sky Survey]



Distribution of Bowhead Whales  
[L.T. Quackenbush et al, *Arctic* 2010]



# KDE Definition



Each point in dataset contributes a *kernel*

Kernel: localized Gaussian “bump”

Kernels summed up to form estimate

Mixture of N Gaussians: N is the dataset size

$$f(x) = \frac{1}{n} \sum_{x_i \in \text{Data}} K(x - x_i)$$


# Problem: KDE does not scale

$$f(x) = \frac{1}{n} \sum_{x_i \in \text{Data}} K(x - x_i)$$

$O(n)$  to compute single density  $f(x)$

$O(n^2)$  to compute all densities in data

2 hours to compute on 1M points  
on 2.9Ghz Core i5




Training Data

How can we speed this up?


# Strawman Optimization: Histograms



Training Dataset



Binned Counts



Grid computation

Benefit: Runtime depends on grid size rather than  $N$


Problem: Bin explosion in high dimensions

[Wand, *J. of Computational and Graphical Statics* 1994]

# Stepping Back: What users need


## Anomaly Explanation

```
SELECT flight_mode FROM shuttle_sensors  
WHERE kde(flow,speed) < threshold
```



## Hypothesis Testing


```
SELECT color FROM galaxies  
WHERE kde(x,y,z) < threshold
```




# From Estimation to Classification

```
SELECT flight_mode FROM shuttle_sensors  
WHERE kde(flow,speed) < Threshold
```


Training Data



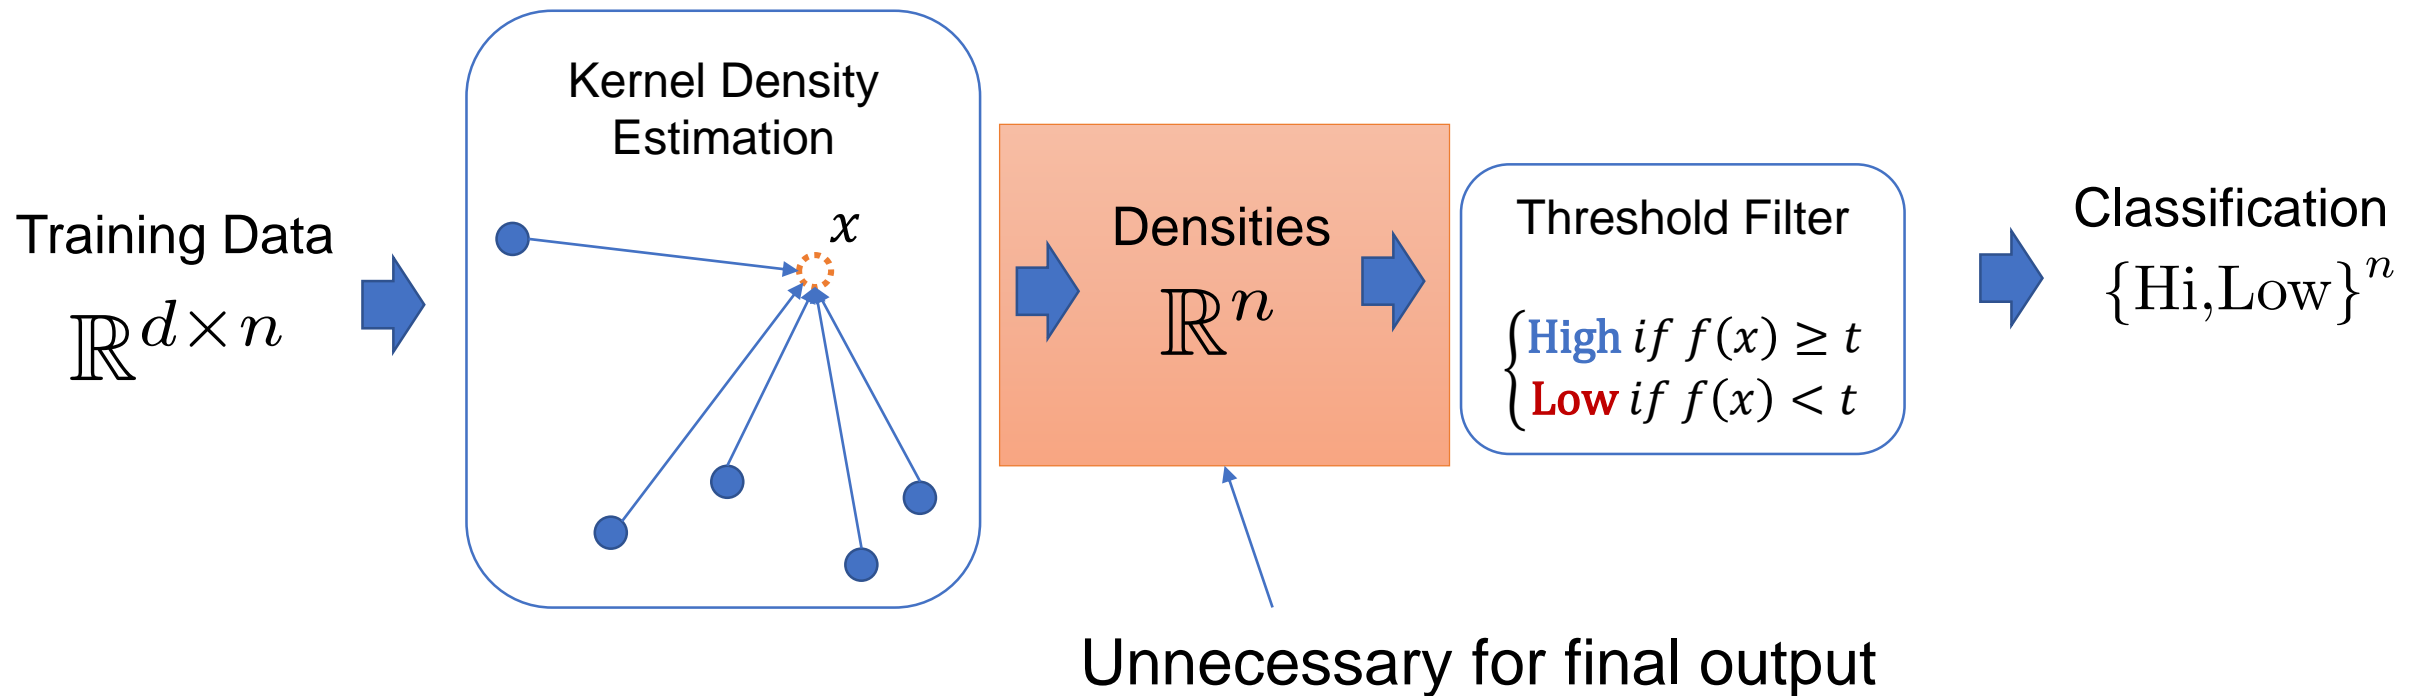
KDE Model



Classification



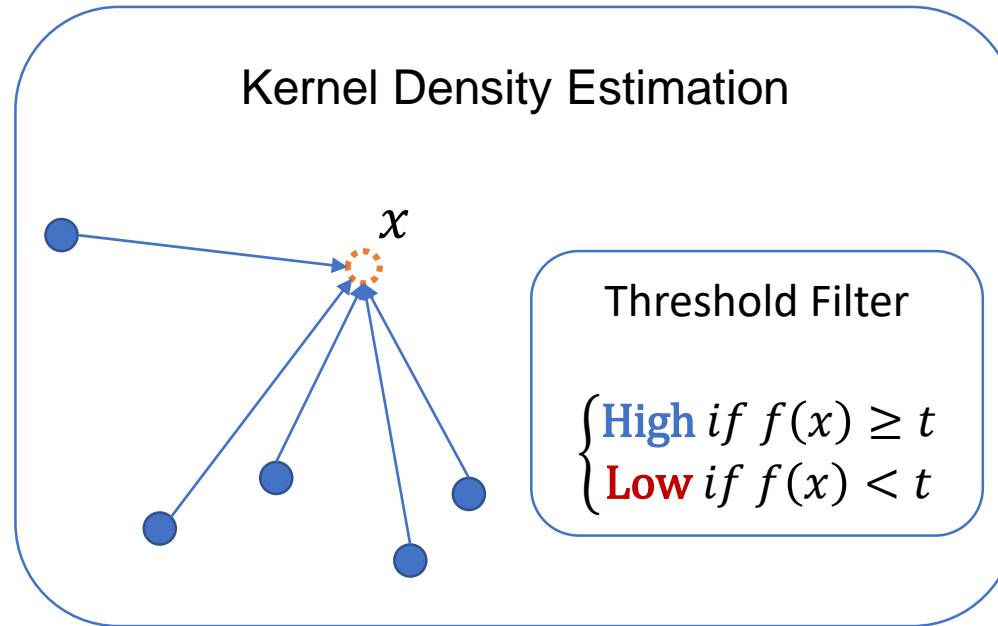
# End to End Query



# End to End Query

Training Data

$$\mathbb{R}^{d \times n}$$



Classification

$$\{\text{Hi}, \text{Low}\}^n$$

# Recap

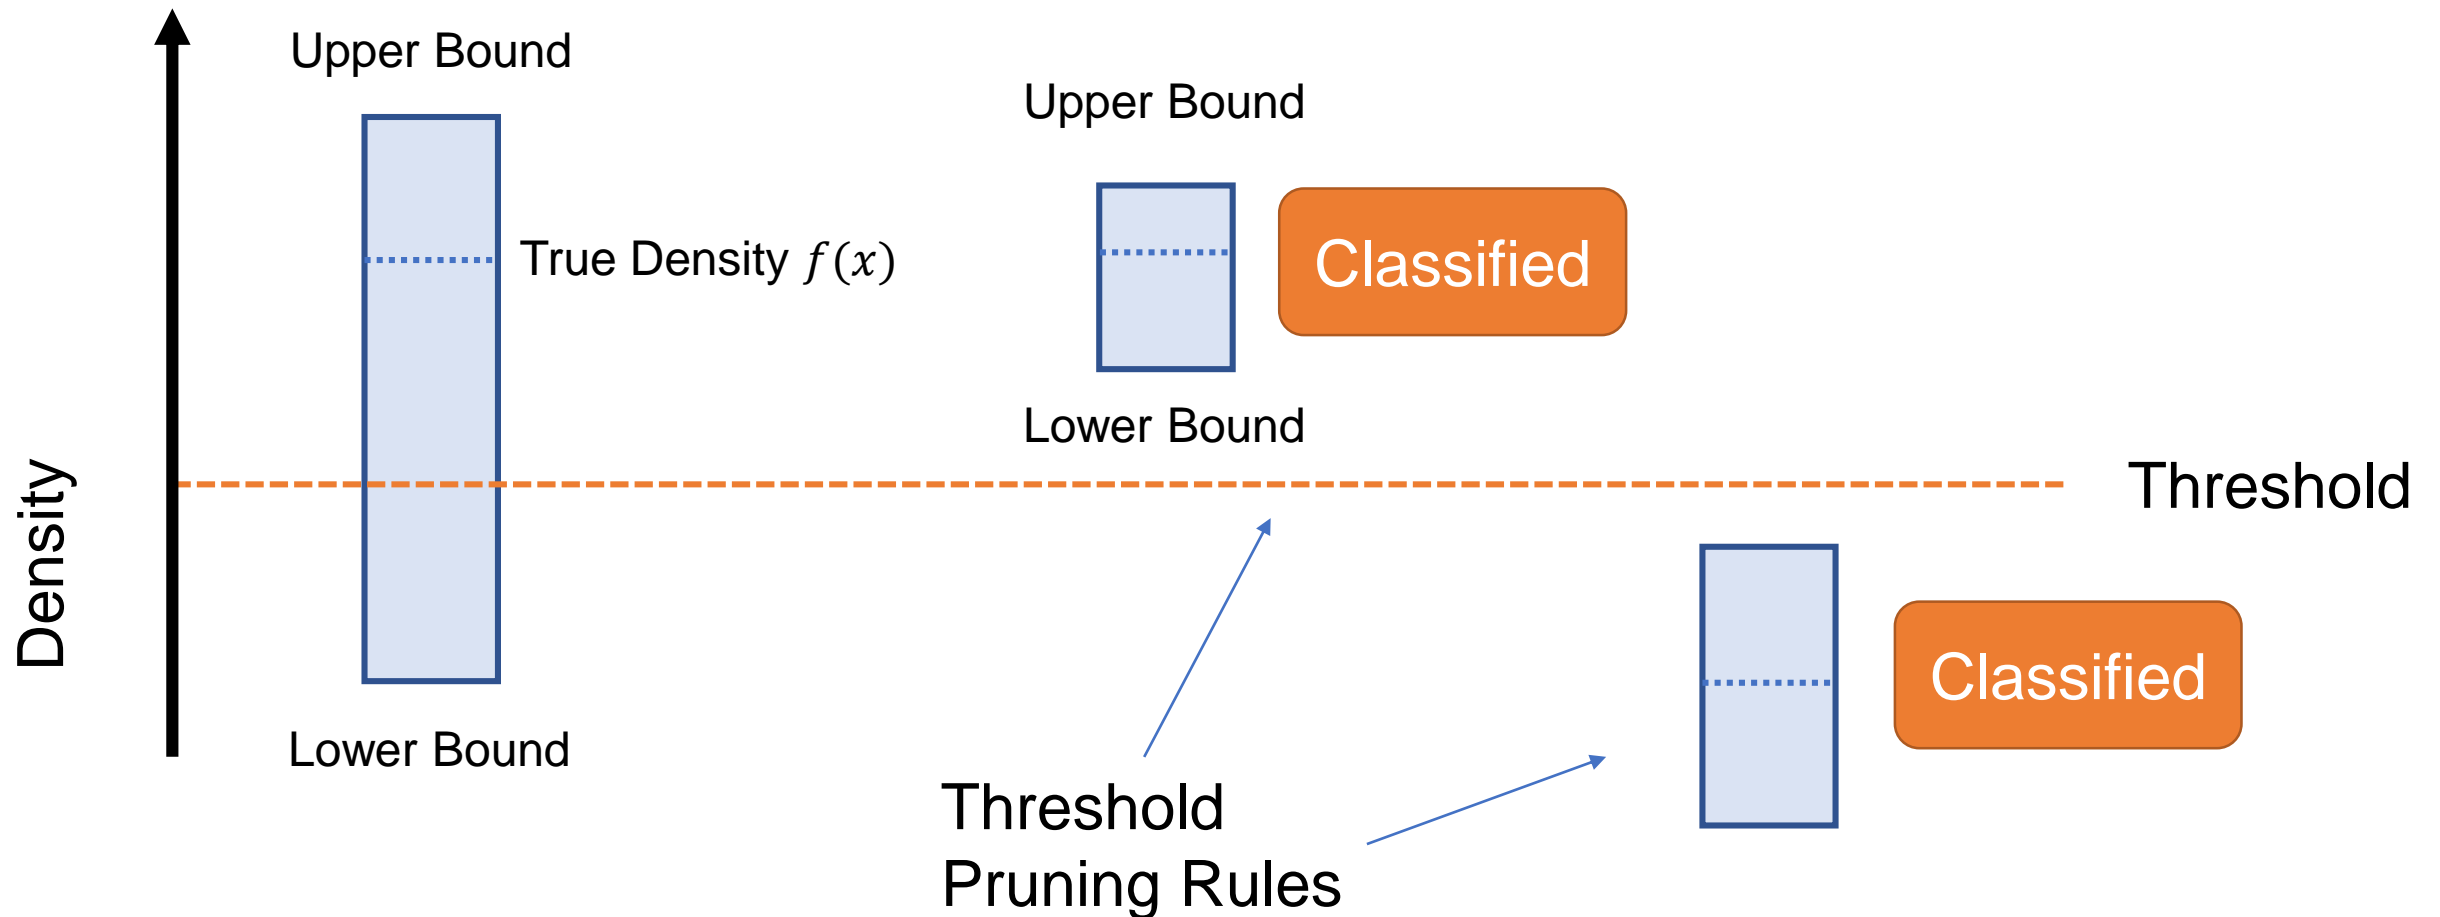
- KDE can model complex distributions
- Problem: KDE scales quadratically with dataset size
- Real Usage: KDE + Predicates = Kernel Density Classification
- Idea: Apply Predicate Pushdown to KDE



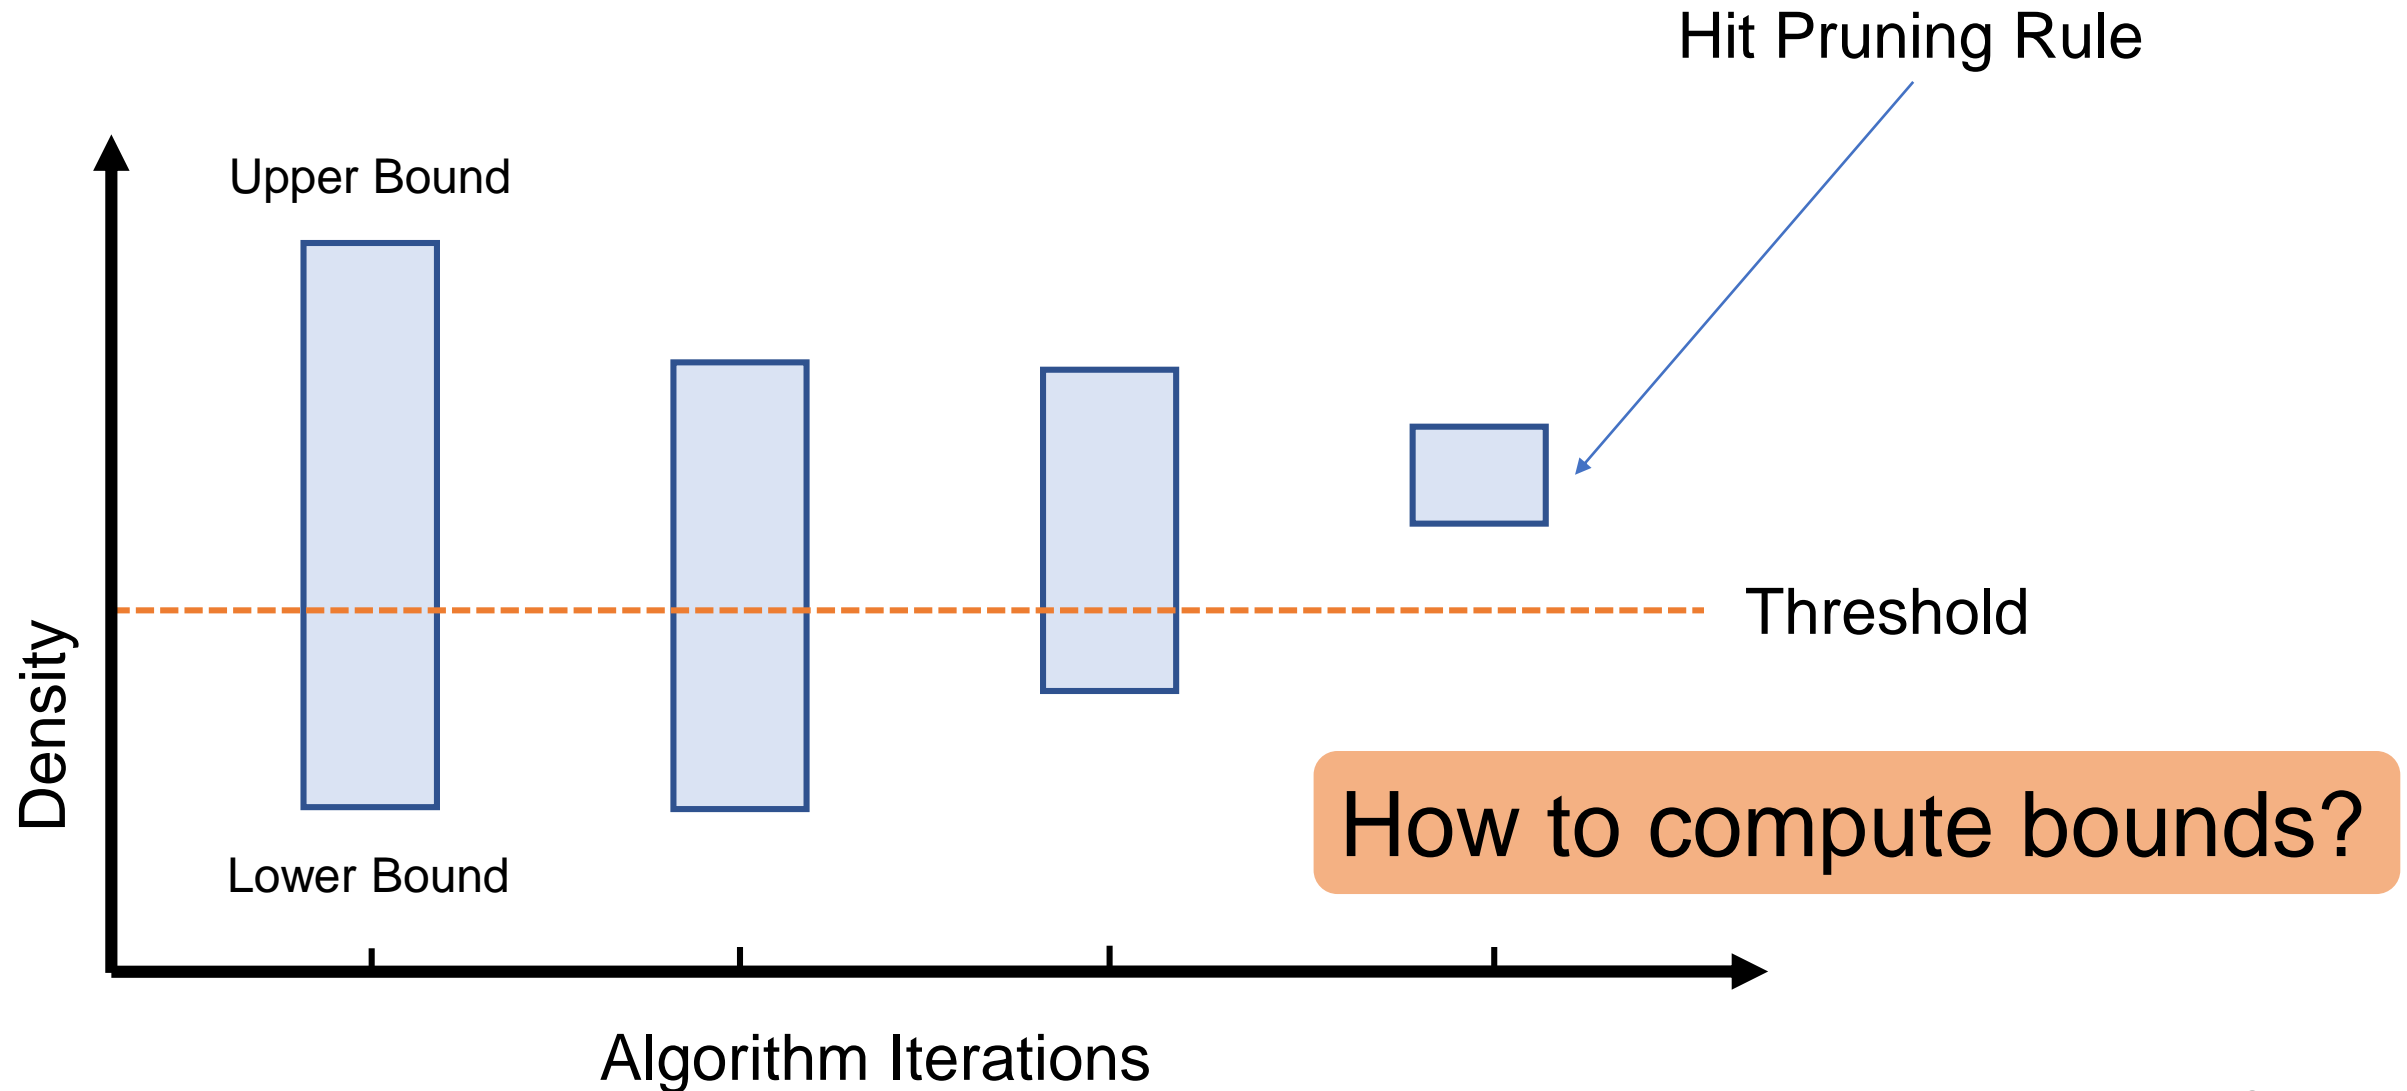
# tkdc Algorithm Overview

1. Pick a threshold
2. Repeat: Calculate bounds on point density
3. Stop when we can make a classification


# Classifying the density based on bounds




# Iterative Refinement



# k-d tree Spatial Indices



Divide N-dimensional space  
1 axis at a time




Nodes for each Region  
Track # of points + bounding box


# Bounding the densities

Given from k-d tree: Bounding Box, # Points Contained


Total contribution from a region can be bounded



Total Contribution




Maximum Contribution



Minimum Contribution


# Iterative Refinement

Initial Estimate




k-d tree root node

Step 1




split

Step 2



split

Step 3



split

**Priority Queue:** Split nodes with largest uncertainty first

# tkdc Algorithm Overview

## 1. Pick a threshold

- User-Specified
- Automatically Inferred

## 2. Calculate bounds on a density

- k-d tree bounding boxes

## 3. Refine the bounds until we can classify


- Priority-queue guided region splitting

# Automatic Threshold Selection

- Probability Densities hard to work with:
  - Unpredictable
  - Huge range of magnitudes
- Good Default: capture a set % of the data

```
SELECT Quantile(kde(A,B), 1%) from shuttle_sensors
```

- Bootstrapping



- Classification for computing thresholds
  - See paper for details



# tkdc Complete Algorithm

- Pick a threshold
  - Inferred given desired % level
- Calculate bounds on a density
  - k-d tree bounding boxes
- Refine the bounds until we can make classification
  - Priority-queue guided region splitting

# Theorem: Expected Runtime

$n$  number of training points

$d$  dimensionality of data

$$\text{Runtime} = O\left(n^{\frac{d-1}{d}}\right)$$

$$\text{Naive} = O(n)$$

100 million data points, 2-dimensions  $\frac{100M}{(100M)^{\frac{1}{2}}} \approx 10,000x$

100 million data points, 8-dimensions  $\frac{100M}{(100M)^{\frac{7}{8}}} \approx 10x$

# Runtime in practice: Experimental Setup

Single Threaded, In-memory



Total Time = Training Time + Threshold Estimation + Classify All

Threshold = 1% classification rate



Baselines:

- simple: naïve for loop over all points
- kdtree: k-d tree approximate density estimation, no threshold
- radial: iterates through points, pruning > certain radius


# KDE Performance Improvement



# Threshold Pruning Contribution



# tkdc scales well with dataset size




# Conclusion

KDE:  
Powerful & Expensive

Real Queries:  
MacroBase

Systems Techniques:

```
SELECT flight_mode FROM shuttle_sensors  
WHERE kde(flow, speed) < Threshold
```



Predicate Pushdown, k-d tree indices:

<https://github.com/stanford-futuredata/tKDC>

1000x, Asymptotic Speedups